

Anscombe's Quartet: Visualization & Regression

EA30 Reflection

October 8, 2025

This handout explores Anscombe's Quartet, a famous dataset demonstrating why data visualization is essential before fitting statistical models. Through guided exercises, students will discover how four datasets with identical summary statistics can have dramatically different patterns.

Introduction: Why Visualization Matters

Anscombe's Quartet is a famous dataset created by statistician Francis Anscombe in 1973 to demonstrate the importance of visualizing data before analyzing it. Four datasets have nearly identical summary statistics (mean, variance, correlation, regression line) but look completely different when graphed!

Key Question to Consider: Can we trust summary statistics alone?

Setup: Loading Required Packages

We need several R packages for data manipulation and visualization: `ggplot2` for creating visualizations, `dplyr` for data manipulation, and `tidyR` for reshaping data.

Data Preparation

The Anscombe dataset comes in "wide" format with separate columns for each set (x_1-y_1 , x_2-y_2 , etc.). We'll transform it to "long" format for easier analysis. We can create 4 dataframes, each with different predictor and response variables, then we can combine them into a single dataframe, which I called `df`, using `rbind()`.

```
# Creating a dataframe with one set
d1 <- data.frame(x = anscombe$x1, y = anscombe$y1, set = "1")
```

Optional: Using tidyR for Data Reshaping

THE METHOD ABOVE works well, but there's a more elegant way using `tidyR`! The `pivot_longer()` function can transform wide data to long format in fewer steps.

Concept: Instead of manually creating four separate data frames, we can:

Learning Goal: Understand why data visualization is essential before fitting statistical models.

Help Topics:

?library
?install.packages

Pseudocode:

1. Load `anscombe` dataset
2. Extract each x-y pair
3. Add set identifiers
4. Combine with `rbind()`

Help Topics:

?anscombe
?data.frame
?rbind

Advanced Technique: The `tidyR` package provides powerful functions for reshaping data.

Help Topics:

?pivot_longer
?separate
?mutate

1. Add a row ID to track observations
2. Pivot all x columns into one column and all y columns into another
3. Separate the set number from the variable name
4. Clean up the result

Here's what the `tidy` approach looks like:

```
# Alternative approach using tidy
df_tidy <- anscombe %>%
  mutate(obs = row_number()) %>%
  pivot_longer(
    cols = -obs,
    names_to = c(".value", "set"),
    names_pattern = "(.)(.)"
  ) %>%
  select(-obs)
```

Exercise 1b (Optional): Comparing Approaches

FOR ADVANCED STUDENTS, try implementing the `tidy` approach:

1. Type `View(anscombe)` to see the original wide format
2. Implement the `tidy` code above
3. Use `identical(df, df_tidy)` to verify both methods produce the same result
4. Which approach do you find more intuitive? Why?

Key Insight: `pivot_longer()` with `names_pattern` recognizes that "`x1`" means "x variable, set 1" and automatically separates them. The `.value` placeholder tells it to create separate columns for x and y.

Exercise 1: Exploring the Data Structure

BEFORE MOVING ON, answer these questions:

1. How many total observations are in the combined `df` dataset?
(Hint: use `nrow(df)`)
2. What are the column names? (Hint: use `names(df)`)
3. Look at the first few rows of dataset 1. What do you notice? (Hint: use `head(subset(df, set == "1"))`)

Bonus: If you completed Exercise 1b above using `tidyR`, compare the structure of both datasets using `str(df)` and `str(df_tidy)`.

Fitting Linear Models

We fit a separate linear regression model ($y \sim x$) to each of the four datasets and extract key statistics: slope, intercept, and R^2 .

Exercise 2: Understanding Model Output

PRACTICE INTERPRETING regression output:

1. Use `summary(fit1)` to see the detailed output for dataset 1. What is the p-value for the slope coefficient?
2. What does the slope tell us about the relationship between x and y ?
3. Try running `plot(fit1)` – what diagnostic plots appear?

Basic Visualization

WHAT IS `ggplot2` – it's a way to visualize data in R that provides a great deal of flexibility.

Pseudocode:

1. Filter data for each set
2. Fit model with `lm()`
3. Extract coefficients
4. Extract R^2 values
5. Create summary table

Help Topics:

?lm
?coef
?summary.lm
?subset

Help Topics:

?ggplot
?geom_point
?geom_smooth
?facet_wrap

```
# Load ggplot2 package
library(ggplot2)


# Create a small example dataset
demo <- data.frame(
  x = 1:10,                      # X-axis values
  y = c(2, 5, 4, 6, 8, 7, 9, 12, 10, 13) # Y-axis values
)

# Create a basic scatter plot
ggplot(demo, aes(x = x, y = y)) +  # Map x and y variables
  geom_point()                    # Add points to the plot
  color = "blue",                 # Point color
  size = 3                        # Point size
) +
  labs()                          # Add labels
```

```

title = "Simple Scatter Plot", # Plot title
x = "X-axis Label",           # X-axis label
y = "Y-axis Label"           # Y-axis label
) +
theme_minimal()             # Apply a minimal theme for clean look

```



Simple Scatter Plot Example


Formatting Options Explained

- **aes(x = x, y = y)**: maps dataset columns to the plot axes.
- **geom_point()**: adds points; can set color, size, shape, etc.
- **labs()**: sets labels like title, subtitle, x-axis, y-axis, and caption.
- **theme_minimal()**: changes plot appearance. Alternatives include `theme_bw()`, `theme_classic()`, etc.
- **Layering with +**: ggplot2 uses + to combine layers like points, lines, labels, and themes.

Notes for EA30

- Each element (`geom_point`, `labs`, `theme_minimal`) can be customized independently.
- You can experiment with colors, sizes, shapes, and themes to see how the plot changes.
- `ggplot2` encourages a "grammar of graphics" approach: map data first (`aes`), then add layers and formatting.

Now we create a 2×2 grid of scatterplots with regression lines to visualize all four datasets. The plot uses `facet_wrap()` to create separate panels for each dataset.



Exercise 3: Visual Interpretation

LOOK CAREFULLY at the four plots and answer:

Figure 1: Scatterplots of all four Anscombe datasets with fitted regression lines. Notice the dramatic differences in patterns despite identical statistical summaries.

1. Which dataset appears to have a truly linear relationship?
2. Which dataset has a curved (non-linear) pattern?
3. Which datasets have outliers? How many in each?
4. Based on visualizations alone, which datasets would be appropriate for linear regression? Why?

Summary Statistics

Notice: All four datasets have nearly identical statistics! Slope ~ 0.50 , Intercept ~ 3.00 , $R^2 \sim 0.67$

set	slope	intercept	r2
1	0.5	3.000	0.667
2	0.5	3.001	0.666
3	0.5	3.002	0.666
4	0.5	3.002	0.667

Table 1: Regression statistics for all four datasets

Exercise 4: Comparing Statistics

CALCULATE ADDITIONAL summary statistics:

1. Calculate the mean of x for each dataset. Are they the same?
(Hint: `tapply(dfx, dfset, mean)`)
2. Calculate the mean of y for each dataset using the same approach
3. Calculate the correlation for each dataset

Question: If the means, correlations, and regression lines are all nearly identical, why do the datasets look so different?

Annotated Visualization

Exercise 5: Customizing Visualization

PRACTICE MODIFYING the plot:

1. Change the point color from "steelblue" to another color
2. Modify the point size – try making them bigger or smaller
3. Change the title to something creative
4. Try `theme_bw()` or `theme_classic()` instead


Critical Observations:

Set 1: Linear (appropriate)

Set 2: Non-linear (curved)

Set 3: Linear + outlier

Set 4: Outlier creates false correlation



Residual Analysis


Residuals help us diagnose whether our linear model is appropriate. Good models have randomly scattered residuals around zero with no clear patterns.

Exercise 6: Interpreting Residual Plots

STUDY THE RESIDUAL PLOTS and answer:

1. Which dataset has residuals that appear randomly scattered? (Good!)
2. Which dataset shows a clear pattern in the residuals? (Linear model inappropriate)
3. Which datasets have extreme residuals (very far from zero)?
4. What do these plots reveal that R^2 did not?

Figure 2: Complete visualization with regression equations displayed on each panel. Despite identical equations, the datasets reveal very different underlying patterns.
Key Concept: Residuals are the differences between observed and predicted y values. Random scatter = good model fit!



Challenge: Impact of Outliers

Exercise 7: Removing Outliers

For DATASET 3 (which has one outlier):

1. Create a new dataset excluding the outlier point ($x=13$)
2. Fit a new regression model to this filtered data
3. Compare the new slope and R^2 to the original
4. Create a scatterplot with the new regression line

Starter hints:

```
df3_no_outlier <- subset(df, set == "3" & x != ?)
fit3_no_outlier <- lm(y ~ x, data = df3_no_outlier)
```

Question: How much did removing one point change your results? What does this tell you about the influence of outliers?

Figure 3: Residual plots reveal problems invisible in summary statistics. Set 1 shows random scatter (good), Set 2 shows clear curvature (bad), and Sets 3 & 4 show extreme outliers.

Final Reflection

SYNTHESIZE WHAT YOU'VE LEARNED:

1. **Main Lesson:** In your own words, what is the key message of Anscombe's Quartet?
2. **Real-World Application:** Think of a situation where someone might make a mistake by looking only at statistics without visualizing data.
3. **Best Practices:** What steps should you ALWAYS take before trusting a linear regression model?
4. **Going Forward:** What other diagnostic tools might you use when analyzing real data?

Key Takeaways

- Always visualize your data before analysis
- Summary statistics can hide important patterns
- Outliers can dramatically affect regression results
- Check model assumptions, don't just trust R² values
- Residual plots help diagnose model problems
- Different data patterns require different analytical approaches

Additional Resources

- **R Documentation:** Use `?function_name` in console
- **ggplot2 Cheat Sheet:** <https://rstudio.github.io/cheatsheets/>
- **Datasaurus Dozen:** A modern extension with even more diverse patterns